

Email Spam Detection

In [16]:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
```

In [2]:

```
df = pd.read_csv("emails.csv")
```

In [3]:

```
df.head()
```

Out[3]:

| | Email No. | the | to | ect | and | for | of | a | you | hou | ... | connevey | jay | valued | lay | infrastructure | militar |
|---|-----------|-----|----|-----|-----|-----|----|-----|-----|-----|-----|----------|-----|--------|-----|----------------|---------|
| 0 | Email 1 | 0 | 0 | 1 | 0 | 0 | 0 | 2 | 0 | 0 | ... | 0 | 0 | 0 | 0 | | 0 |
| 1 | Email 2 | 8 | 13 | 24 | 6 | 6 | 2 | 102 | 1 | 27 | ... | 0 | 0 | 0 | 0 | | 0 |
| 2 | Email 3 | 0 | 0 | 1 | 0 | 0 | 0 | 8 | 0 | 0 | ... | 0 | 0 | 0 | 0 | | 0 |
| 3 | Email 4 | 0 | 5 | 22 | 0 | 5 | 1 | 51 | 2 | 10 | ... | 0 | 0 | 0 | 0 | | 0 |
| 4 | Email 5 | 7 | 6 | 17 | 1 | 5 | 2 | 57 | 0 | 9 | ... | 0 | 0 | 0 | 0 | | 0 |

5 rows × 3002 columns

In [4]:

```
df.isnull().sum()
```

Out[4]:

```
Email No.      0
the            0
to            0
ect           0
and           0
..
military      0
allowing     0
ff           0
dry          0
Prediction   0
Length: 3002, dtype: int64
```

In [5]:

```
X = df.iloc[:,1:3001]
X
```

Out[5]:

| | the | to | ect | and | for | of | a | you | hou | in | ... | enhancements | connevey | jay | valued | lay | infra |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|--------------|----------|-----|--------|-----|-------|
| 0 | 0 | 0 | 1 | 0 | 0 | 0 | 2 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 8 | 13 | 24 | 6 | 6 | 2 | 102 | 1 | 27 | 18 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 1 | 0 | 0 | 0 | 8 | 0 | 0 | 4 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 3 | 0 | 5 | 22 | 0 | 5 | 1 | 51 | 2 | 10 | 1 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 4 | 7 | 6 | 17 | 1 | 5 | 2 | 57 | 0 | 9 | 3 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 5167 | 2 | 2 | 2 | 3 | 0 | 0 | 32 | 0 | 0 | 5 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 5168 | 35 | 27 | 11 | 2 | 6 | 5 | 151 | 4 | 3 | 23 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 5169 | 0 | 0 | 1 | 1 | 0 | 0 | 11 | 0 | 0 | 1 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 5170 | 2 | 7 | 1 | 0 | 2 | 1 | 28 | 2 | 0 | 8 | ... | 0 | 0 | 0 | 0 | 0 | 0 |
| 5171 | 22 | 24 | 5 | 1 | 6 | 5 | 148 | 8 | 2 | 23 | ... | 0 | 0 | 0 | 0 | 0 | 0 |

5172 rows × 3000 columns

In [6]:

```
Y = df.iloc[:, -1].values
Y
```

Out[6]:

```
array([0, 0, 0, ..., 1, 1, 0])
```

In [7]:

```
train_x, test_x, train_y, test_y = train_test_split(X,Y,test_size = 0.25)
```

In [8]:

```
svc = SVC(C=1.0,kernel='rbf',gamma='auto')
# C here is the regularization parameter. Here, L2 penalty is used(default). It is the inverse of the regularization strength.
# As C increases, model overfits.
# Kernel here is the radial basis function kernel.
# gamma (only used for rbf kernel) : As gamma increases, model overfits.
svc.fit(train_x,train_y)
y_pred2 = svc.predict(test_x)
```

In [9]:

```
print("Accuracy Score for SVC : ", accuracy_score(y_pred2,test_y))
```

```
Accuracy Score for SVC : 0.9040989945862336
```

In [10]:

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.2, random_state=42)
```

In [11]:

```
knn = KNeighborsClassifier(n_neighbors=7)
```

In [12]:

```
knn.fit(X_train, y_train)
```

Out[12]:

```
KNeighborsClassifier(n_neighbors=7)
```

In [13]:

```
print(knn.predict(X_test))
```

```
[0 0 1 ... 0 1 0]
```

In [14]:

```
print(knn.score(X_test, y_test))
```

```
0.8685990338164251
```

In [17]:

```
val = df['Prediction'].value_counts()
val.plot(kind='bar', color='gray')
plt.xlabel('')
plt.ylabel('Quantities')
plt.title('Quantities of Spam Emails')

plt.xticks([0,1],['Not Spam','Spam'])
plt.show()
```

