

Implement K-Means clustering/ hierarchical clustering on sales_data_sample.csv dataset. Determine the number of clusters using the elbow method.

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: df = pd.read_csv('sales_data_sample.csv', encoding='unicode_escape')
```

```
In [6]: df.columns
```

```
Out[6]: Index(['ORDERNUMBER', 'QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER',  
              'SALES', 'ORDERDATE', 'STATUS', 'QTR_ID', 'MONTH_ID', 'YEAR_ID',  
              'PRODUCTLINE', 'MSRP', 'PRODUCTCODE', 'CUSTOMERNAME', 'PHONE',  
              'ADDRESSLINE1', 'ADDRESSLINE2', 'CITY', 'STATE', 'POSTALCODE',  
              'COUNTRY', 'TERRITORY', 'CONTACTLASTNAME', 'CONTACTFIRSTNAME',  
              'DEALSIZE'],  
             dtype='object')
```

```
In [3]: df.head
```

```
Out[3]: <bound method NDFrame.head of
ORDERLINENUMBER  SALES  \
0      10107      30      95.70      2  2871.00
1      10121      34      81.35      5  2765.90
2      10134      41      94.74      2  3884.34
3      10145      45      83.26      6  3746.70
4      10159      49     100.00     14  5205.27
...      ...      ...      ...      ...      ...
2818     10350      20     100.00     15  2244.40
2819     10373      29     100.00      1  3978.51
2820     10386      43     100.00      4  5417.57
2821     10397      34      62.24      1  2116.16
2822     10414      47      65.52      9  3079.44

ORDERDATE  STATUS  QTR_ID  MONTH_ID  YEAR_ID  ...  \
0      2/24/2003 0:00  Shipped      1      2     2003  ...
1      5/7/2003 0:00  Shipped      2      5     2003  ...
2      7/1/2003 0:00  Shipped      3      7     2003  ...
3      8/25/2003 0:00  Shipped      3      8     2003  ...
4     10/10/2003 0:00  Shipped      4     10     2003  ...
...      ...      ...      ...      ...      ...  ...
2818    12/2/2004 0:00  Shipped      4     12     2004  ...
2819    1/31/2005 0:00  Shipped      1      1     2005  ...
2820     3/1/2005 0:00  Resolved      1      3     2005  ...
2821    3/28/2005 0:00  Shipped      1      3     2005  ...
2822     5/6/2005 0:00  On Hold      2      5     2005  ...

ADDRESSLINE1  ADDRESSLINE2  CITY  STATE  \
0      897 Long Airport Avenue      NaN      NYC      NY
1           59 rue de l'Abbaye      NaN      Reims      NaN
2      27 rue du Colonel Pierre Avia      NaN      Paris      NaN
3           78934 Hillside Dr.      NaN      Pasadena      CA
4           7734 Strong St.      NaN  San Francisco      CA
...      ...      ...      ...      ...
2818           C/ Moralzarzal, 86      NaN      Madrid      NaN
2819           Torikatu 38      NaN      Oulu      NaN
2820           C/ Moralzarzal, 86      NaN      Madrid      NaN
2821      1 rue Alsace-Lorraine      NaN      Toulouse      NaN
2822           8616 Spinnaker Dr.      NaN      Boston      MA

POSTALCODE  COUNTRY  TERRITORY  CONTACTLASTNAME  CONTACTFIRSTNAME  DEALSIZE
0      10022      USA      NaN      Yu      Kwai      Small
1      51100  France      EMEA      Henriot      Paul      Small
2      75508  France      EMEA      Da Cunha      Daniel      Medium
3      90003      USA      NaN      Young      Julie      Medium
4      NaN      USA      NaN      Brown      Julie      Medium
...      ...      ...      ...      ...      ...
2818     28034  Spain      EMEA      Freyre      Diego      Small
2819     90110  Finland      EMEA      Koskitalo      Pirkko      Medium
2820     28034  Spain      EMEA      Freyre      Diego      Medium
2821     31000  France      EMEA      Roulet      Annette      Small
2822     51003      USA      NaN      Yoshido      Juri      Medium
```

```
[2823 rows x 25 columns]>
```

In [4]: df.info

```
Out[4]: <bound method DataFrame.info of
ORDERLINENUMBER  SALES  \
0      10107      30      95.70      2  2871.00
1      10121      34      81.35      5  2765.90
2      10134      41      94.74      2  3884.34
3      10145      45      83.26      6  3746.70
4      10159      49     100.00     14  5205.27
...      ...      ...      ...      ...      ...
2818     10350      20     100.00     15  2244.40
2819     10373      29     100.00      1  3978.51
2820     10386      43     100.00      4  5417.57
2821     10397      34      62.24      1  2116.16
2822     10414      47      65.52      9  3079.44

      ORDERDATE  STATUS  QTR_ID  MONTH_ID  YEAR_ID  ...  \
0      2/24/2003 0:00  Shipped      1         2     2003  ...
1      5/7/2003 0:00  Shipped      2         5     2003  ...
2      7/1/2003 0:00  Shipped      3         7     2003  ...
3      8/25/2003 0:00  Shipped      3         8     2003  ...
4     10/10/2003 0:00  Shipped      4        10     2003  ...
...      ...      ...      ...      ...      ...      ...
2818    12/2/2004 0:00  Shipped      4        12     2004  ...
2819    1/31/2005 0:00  Shipped      1         1     2005  ...
2820     3/1/2005 0:00  Resolved      1         3     2005  ...
2821    3/28/2005 0:00  Shipped      1         3     2005  ...
2822     5/6/2005 0:00  On Hold      2         5     2005  ...

      ADDRESSLINE1 ADDRESSLINE2      CITY STATE  \
0      897 Long Airport Avenue      NaN      NYC  NY
1           59 rue de l'Abbaye      NaN      Reims  NaN
2      27 rue du Colonel Pierre Avia      NaN      Paris  NaN
3           78934 Hillside Dr.      NaN      Pasadena  CA
4           7734 Strong St.      NaN  San Francisco  CA
...      ...      ...      ...      ...
2818      C/ Moralzarzal, 86      NaN      Madrid  NaN
2819           Torikatu 38      NaN      Oulu  NaN
2820      C/ Moralzarzal, 86      NaN      Madrid  NaN
2821      1 rue Alsace-Lorraine      NaN      Toulouse  NaN
2822      8616 Spinnaker Dr.      NaN      Boston  MA

      POSTALCODE  COUNTRY  TERRITORY  CONTACTLASTNAME  CONTACTFIRSTNAME  DEALSIZE
0      10022      USA      NaN      Yu      Kwai      Small
1      51100  France      EMEA      Henriot      Paul      Small
2      75508  France      EMEA      Da Cunha      Daniel      Medium
3      90003      USA      NaN      Young      Julie      Medium
4      NaN      USA      NaN      Brown      Julie      Medium
...      ...      ...      ...      ...      ...
2818     28034  Spain      EMEA      Freyre      Diego      Small
2819     90110  Finland      EMEA      Koskitalo      Pirkko      Medium
2820     28034  Spain      EMEA      Freyre      Diego      Medium
2821     31000  France      EMEA      Roulet      Annette      Small
2822     51003      USA      NaN      Yoshido      Juri      Medium
```

[2823 rows x 25 columns]>

```
In [7]: #Columns to Remove
to_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATE', 'POSTALCODE', 'PHONE']
df = df.drop(to_drop, axis=1)
```

```
In [8]: #Check for null values
df.isnull().sum()
```

```
Out[8]: ORDERNUMBER          0
QUANTITYORDERED          0
PRICEEACH                0
ORDERLINENUMBER          0
SALES                    0
ORDERDATE                0
STATUS                   0
QTR_ID                   0
MONTH_ID                 0
YEAR_ID                  0
PRODUCTLINE              0
MSRP                     0
PRODUCTCODE              0
CUSTOMERNAME             0
CITY                     0
COUNTRY                  0
TERRITORY                1074
CONTACTLASTNAME          0
CONTACTFIRSTNAME         0
DEALSIZE                  0
dtype: int64
```

```
In [9]: df.dtypes
```

```
Out[9]: ORDERNUMBER          int64
QUANTITYORDERED          int64
PRICEEACH                float64
ORDERLINENUMBER          int64
SALES                    float64
ORDERDATE                object
STATUS                   object
QTR_ID                   int64
MONTH_ID                 int64
YEAR_ID                  int64
PRODUCTLINE              object
MSRP                     int64
PRODUCTCODE              object
CUSTOMERNAME             object
CITY                     object
COUNTRY                  object
TERRITORY                object
CONTACTLASTNAME          object
CONTACTFIRSTNAME         object
DEALSIZE                  object
dtype: object
```

```
In [10]: #ORDERDATE Should be in date time
df['ORDERDATE'] = pd.to_datetime(df['ORDERDATE'])
```

```
In [11]: #We need to create some features in order to create cluserers
#Recency: Number of days between customer's latest order and today's date
#Frequency : Number of purchases by the customers
#MonetaryValue : Revenue generated by the customers
```

```
import datetime as dt
snapshot_date = df['ORDERDATE'].max() + dt.timedelta(days = 1)
df_RFM = df.groupby(['CUSTOMERNAME']).agg({
    'ORDERDATE' : lambda x : (snapshot_date - x.max()).days,
    'ORDERNUMBER' : 'count',
    'SALES' : 'sum'
})

#Rename the columns
df_RFM.rename(columns = {
    'ORDERDATE' : 'Recency',
    'ORDERNUMBER' : 'Frequency',
    'SALES' : 'MonetaryValue'
}, inplace=True)
```

```
In [12]: df_RFM.head()
```

```
Out[12]:
```

	Recency	Frequency	MonetaryValue
CUSTOMERNAME			
AV Stores, Co.	196	51	157807.81
Alpha Cognac	65	20	70488.44
Amica Models & Co.	265	26	94117.26
Anna's Decorations, Ltd	84	46	153996.13
Atelier graphique	188	7	24179.96

```
In [13]: # Divide into segments
# We create 4 quartile ranges
df_RFM['M'] = pd.qcut(df_RFM['MonetaryValue'], q = 4, labels = range(1,5))
df_RFM['R'] = pd.qcut(df_RFM['Recency'], q = 4, labels = list(range(4,0,-1)))
df_RFM['F'] = pd.qcut(df_RFM['Frequency'], q = 4, labels = range(1,5))

df_RFM.head()
```

Out[13]:

	Recency	Frequency	MonetaryValue	M	R	F
CUSTOMERNAME						
AV Stores, Co.	196	51	157807.81	4	2	4
Alpha Cognac	65	20	70488.44	2	4	2
Amica Models & Co.	265	26	94117.26	3	1	2
Anna's Decorations, Ltd	84	46	153996.13	4	3	4
Atelier graphique	188	7	24179.96	1	2	1

```
In [14]: #Create another column for RFM score
df_RFM['RFM_Score'] = df_RFM[['R', 'M', 'F']].sum(axis=1)
df_RFM.head()
```

Out[14]:

	Recency	Frequency	MonetaryValue	M	R	F	RFM_Score
CUSTOMERNAME							
AV Stores, Co.	196	51	157807.81	4	2	4	10
Alpha Cognac	65	20	70488.44	2	4	2	8
Amica Models & Co.	265	26	94117.26	3	1	2	6
Anna's Decorations, Ltd	84	46	153996.13	4	3	4	11
Atelier graphique	188	7	24179.96	1	2	1	4

We create levels for our Customers
 RFM Score > 10 : High Value Customers
 RFM Score < 10 and RFM Score >= 6 : Mid Value Customers
 RFM Score < 6 : Low Value Customers

```
In [15]: def rfm_level(df):
    if bool(df['RFM_Score'] >= 10):
        return 'High Value Customer'

    elif bool(df['RFM_Score'] < 10) and bool(df['RFM_Score'] >= 6):
        return 'Mid Value Customer'
    else:
        return 'Low Value Customer'
df_RFM['RFM_Level'] = df_RFM.apply(rfm_level, axis = 1)
df_RFM.head()
```

Out[15]:

	Recency	Frequency	MonetaryValue	M	R	F	RFM_Score	RFM_Level
CUSTOMERNAME								
AV Stores, Co.	196	51	157807.81	4	2	4	10	High Value Customer
Alpha Cognac	65	20	70488.44	2	4	2	8	Mid Value Customer
Amica Models & Co.	265	26	94117.26	3	1	2	6	Mid Value Customer
Anna's Decorations, Ltd	84	46	153996.13	4	3	4	11	High Value Customer
Atelier graphique	188	7	24179.96	1	2	1	4	Low Value Customer

```
In [16]: # Time to perform KMeans
data = df_RFM[['Recency', 'Frequency', 'MonetaryValue']]
data.head()
```

Out[16]:

	Recency	Frequency	MonetaryValue
CUSTOMERNAME			
AV Stores, Co.	196	51	157807.81
Alpha Cognac	65	20	70488.44
Amica Models & Co.	265	26	94117.26
Anna's Decorations, Ltd	84	46	153996.13
Atelier graphique	188	7	24179.96

```
In [17]: # Our data is skewed we must remove it by performing log transformation
data_log = np.log(data)
data_log.head()
```

Out[17]:

	Recency	Frequency	MonetaryValue
CUSTOMERNAME			
AV Stores, Co.	5.278115	3.931826	11.969133
Alpha Cognac	4.174387	2.995732	11.163204
Amica Models & Co.	5.579730	3.258097	11.452297
Anna's Decorations, Ltd	4.430817	3.828641	11.944683
Atelier graphique	5.236442	1.945910	10.093279

```
In [18]: #Standardization
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(data_log)
data_normalized = scaler.transform(data_log)
data_normalized = pd.DataFrame(data_normalized, index = data_log.index, columns = data_log.columns)
data_normalized.describe().round(2)
```

Out[18]:

	Recency	Frequency	MonetaryValue
count	92.00	92.00	92.00
mean	0.00	-0.00	0.00
std	1.01	1.01	1.01
min	-3.51	-3.67	-3.82
25%	-0.24	-0.41	-0.39
50%	0.37	0.06	-0.04
75%	0.53	0.45	0.52
max	1.12	4.03	3.92

In [22]: *#Fit KMeans and use elbow method to choose the number of clusters*

```
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.cluster import KMeans


sse = {}

for k in range(1, 21):
    kmeans = KMeans(n_clusters = k, random_state = 1)
    kmeans.fit(data_normalized)
    sse[k] = kmeans.inertia_
```

In [26]: `plt.figure(figsize=(10,6))`
`plt.title('The Elbow Method')`

`plt.xlabel('K')`
`plt.ylabel('SSE')`
`plt.style.use('ggplot')`

`sns.pointplot(x=list(sse.keys()), y = list(sse.values()))`
`plt.text(4.5, 60, "Largest Angle", bbox = dict(facecolor = 'lightgreen', alpha`
`plt.show()`



```
In [27]: # 5 number of clusters seems good
kmeans = KMeans(n_clusters=5, random_state=1)
kmeans.fit(data_normalized)
cluster_labels = kmeans.labels_

data_rfm = data.assign(Cluster = cluster_labels)
data_rfm.head()
```

Out[27]:

	Recency	Frequency	MonetaryValue	Cluster
CUSTOMERNAME				
AV Stores, Co.	196	51	157807.81	3
Alpha Cognac	65	20	70488.44	0
Amica Models & Co.	265	26	94117.26	0
Anna's Decorations, Ltd	84	46	153996.13	3
Atelier graphique	188	7	24179.96	2